

MISP

INTELLIGENCE SHARING

MISP Galaxy in 2.5

Explore the knowledge base, assess context, share carefully

CIRCL / MISP Project

https://www.misp-project.org/

Reference snapshot: 30 October 2026

Find context, assess evidence, make an assertion

- Explore the knowledge base beyond actor names.
- Choose the right cluster and attachment scope.
- Interpret relations and matrices carefully.
- Practise discovery, custom knowledge and contribution.

Explore the knowledge base

A large, reusable body of knowledge

- 136 Galaxy definitions and 136 cluster collections.
- 58,528 entries; 56,551 distinct cluster UUIDs.
- 79,372 listed directed relationship edges.

Start with the public catalogue

- Browse https://www.misp-galaxy.org/.
- Follow JSON and original-source references.
- Inspect collection scope, authors, version and status.

Choose a family that answers the question

- ATT&CK: observed techniques, software and activity groups.
- D3FEND and preventive measures: defensive context.
- DISARM: information-manipulation techniques and responses.
- Country, sector, detection and other domain knowledge.

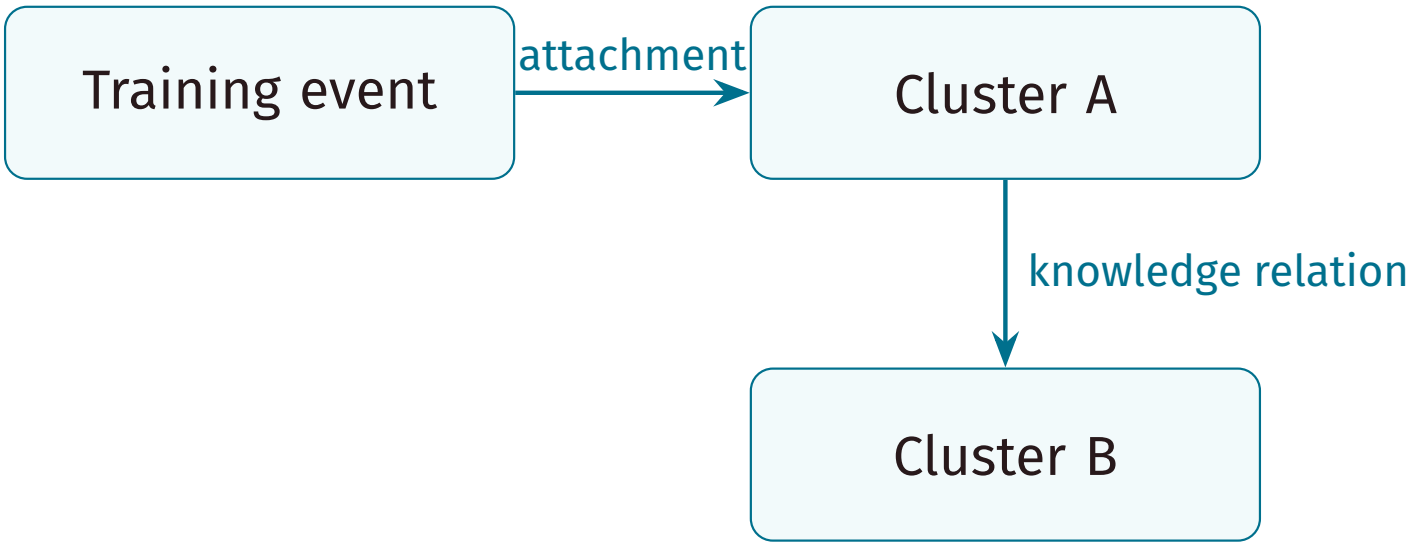
Galaxy, cluster and element

- Galaxy: the family and its presentation metadata.
- Cluster: one named concept with an individual UUID.
- Element: metadata such as synonyms, references or external IDs.

Context is different from incident data

- Taxonomy: classification vocabulary.
- MISP object: grouped typed attributes.
- Galaxy cluster: structured knowledge attached as context.

Three different links



An attachment does not create a cluster-to-cluster edge.

Search for identity, not just a familiar label

- Compare values, synonyms and external identifiers.
- Retain cluster UUID and collection provenance.
- Do not merge identities solely because names overlap.

Check whether the entry is current

- Read original references and collection version.
- Inspect revoked flags, deprecation metadata and historical descriptions.
- Check ATT&CK domain, platform and technique/sub-technique scope.

Apply context in MISP 2.5

Choose the narrowest justified attachment

- Event scope describes the incident or report as a whole.
- Attribute scope can express a narrower observation.
- Verify the actual attachment after saving and exporting.

Use the installed cluster metadata

- Search and inspect the cluster before attaching it.
- MISP maintains a cluster-associated tag.
- Prefer UUID-backed identity over manually constructed labels.

Keep visibility boundaries separate

- Event visibility controls incident access.
- Custom-cluster distribution controls knowledge access.
- Local attachment controls the association scope.
- Publication and sync settings control eligible sharing.

Matrix highlights are observations, not coverage

- Galaxy ordering defines matrix tabs and columns.
- Cluster kill-chain metadata places entries in that layout.
- A highlight does not establish detection or prevention efficacy.

Read graph edges as sourced assertions

- Inspect direction, relationship type and references.
- Review confidence or handling tags where supplied.
- Do not infer causality merely from a graph path.

Custom clusters and forks are supported

- Create permitted custom knowledge with explicit ownership.
- Fork a default cluster for local analysis.
- Record provenance, distribution and publication intent.

Updating files is different from updating the catalogue

- Select and update the intended Git revision.
- Import/update the MISP database through the operator workflow.
- Review changed knowledge and custom forks before sharing.

Explore, validate and contribute

Explore a pinned checkout offline

```
python3 examples/explore.py /tmp/misp-galaxy-lab --stats
python3 examples/explore.py /tmp/misp-galaxy-lab \
  --search T1059 \
  --collection mitre-enterprise-attack-attack-pattern \
  --limit 3
```

Repository structure: family versus collection

```
galaxies/training-knowledge.json
  type, namespace, name, uuid, version, description

clusters/training-knowledge.json
  type, name, uuid, version, source, authors, category
  values: [individual cluster entries]
```

An individual cluster can reference another

```
{
  "value": "Training observation A",
  "uuid": "<individual cluster UUID>",
  "meta": {"synonyms": ["Workshop A"]},
  "related": [{
    "dest-uuid": "<cluster B UUID>",
    "type": "related-to", "tags": []
  }]
}
```

Validate more than JSON syntax

- Check both repository schemas.
- Check identities, matching types and relationship targets.
- Exercise negative cases and inspect original sources.

Workshop: find, assess, attach

- Find a technique by external identifier and read its evidence.
- Compare a synonym search and a domain-specific collection.
- Attach to a disposable event; compare attribute and local scope.
- Inspect matrix/export results and remove training data.

Contribute at the correct source

- Manual collection: edit JSON with evidence and provenance.
- Generated collection: inspect the generator and upstream source.
- Preserve identities and update relevant versions.

Follow the current upstream validation workflow

```
./jq_all_the_things.sh
# Review and commit intended normalised changes.
./validate_all.sh
```

Make knowledge reusable and reviewable

- State source, identity, meaning and uncertainty.
- Keep attachment scope separate from knowledge distribution.
- Include reproducible examples and validation limits.
- Use the README for pinned references and workshop commands.